

```

1 !pip install -q torchvision
2
3 import os, shutil
4 import numpy as np
5 from pathlib import Path
6 import matplotlib.pyplot as plt
7 from PIL import Image
8 from tqdm.auto import tqdm
9
10 import torch
11 import torchvision
12 from torch.utils.data import DataLoader
13 from torchvision.models import resnet50, ResNet50_Weights
14 from torchvision.transforms import transforms
15 from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, ConfusionMatrixDisplay, f1_score
16

```

```

→ _____ 363.4/363.4 MB 2.9 MB/s eta 0:00:00
_____ 13.8/13.8 MB 111.9 MB/s eta 0:00:00
_____ 24.6/24.6 MB 81.1 MB/s eta 0:00:00
_____ 883.7/883.7 kB 52.2 MB/s eta 0:00:00
_____ 664.8/664.8 MB 1.3 MB/s eta 0:00:00
_____ 211.5/211.5 MB 9.9 MB/s eta 0:00:00
_____ 56.3/56.3 MB 39.3 MB/s eta 0:00:00
_____ 127.9/127.9 MB 7.7 MB/s eta 0:00:00
_____ 207.5/207.5 MB 10.2 MB/s eta 0:00:00
_____ 21.1/21.1 MB 102.1 MB/s eta 0:00:00

```

```

1 # Define the transformation pipeline for input images
2 transform = transforms.Compose([
3     transforms.Resize((224,224)), # Resize input image to 224x224 (expected by ResNet)
4     transforms.ToTensor()         # Convert image to PyTorch tensor
5 ])
6
7 # Define a custom feature extractor class using ResNet-50
8 class resnet_feature_extractor(torch.nn.Module):
9     def __init__(self):
10         super(resnet_feature_extractor, self).__init__()
11
12         # Load pretrained ResNet-50 model with default weights
13         self.model = resnet50(weights=ResNet50_Weights.DEFAULT)
14         self.model.eval() # Set model to evaluation mode (disable dropout, batchnorm updates)
15
16         # Freeze all model parameters (no gradient updates)
17         for param in self.model.parameters():
18             param.requires_grad = False
19
20         # Define a forward hook to capture feature maps from intermediate layers
21         def hook(module, input, output):
22             self.features.append(output)
23
24         # Register hooks on the last block of layer2 and layer3
25         self.model.layer2[-1].register_forward_hook(hook)
26         self.model.layer3[-1].register_forward_hook(hook)
27
28     def forward(self, x):
29         self.features = [] # Clear any previously stored features
30
31         # Forward pass without gradient computation
32         with torch.no_grad():
33             _ = self.model(x) # Run input through the full ResNet model
34
35         # Apply 2D average pooling to smooth feature maps
36         self.avg = torch.nn.AvgPool2d(kernel_size=3, stride=1)
37
38         # Determine the spatial size of the first captured feature map
39         fmap_size = self.features[0].shape[-2]
40
41         # Resize all feature maps to have the same spatial size using adaptive pooling
42         self.resize = torch.nn.AdaptiveAvgPool2d(fmap_size)
43
44         # Apply pooling and resizing to all captured feature maps
45         resized_maps = [self.resize(self.avg(fmap)) for fmap in self.features]
46
47         # Concatenate feature maps along the channel dimension
48         patch = torch.cat(resized_maps, dim=1) # shape: (B, C_total, H, W)

```

```

49
50     # Flatten each spatial location's concatenated features to form a patch matrix
51     patch = patch.reshape(patch.shape[1], -1).T # shape: (H*W, C_total)
52
53     return patch # Return patch matrix of shape (num_patches, total_channels)
54

```

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3

```

Mounted at /content/drive

```

1 # Initialize an empty list to store extracted features (memory bank)
2 memory_bank = []
3
4 # Instantiate the ResNet-based feature extractor and move it to GPU
5 backbone = resnet_feature_extractor().cuda()
6
7 # Define the path to the training "good" images in the MVTec zipper dataset
8 folder_path = Path('/content/drive/MyDrive/mvttec_anomaly_detection/metal_nut/train/good')
9
10 # Iterate over all image files in the folder
11 for pth in tqdm(folder_path.iterdir(), leave=False):
12     # Load the image, convert to RGB (ensure no alpha channel), apply transforms
13     with torch.no_grad(): # Disable gradient computation for faster inference
14         data = transform(Image.open(pth).convert("RGB")).cuda().unsqueeze(0) # Shape: (1, 3, 224, 224)
15
16     # Extract features using the backbone
17     features = backbone(data) # Shape: (num_patches, feature_dim) e.g., (784, 391)
18
19     # Store features in CPU memory to save GPU memory
20     memory_bank.append(features.cpu().detach())
21
22 # Print how many images were processed
23 print(len(memory_bank)) # Should equal number of images in the folder
24
25 # Print the shape of the feature map for the first image
26 print(memory_bank[0].shape) # e.g., torch.Size([784, 391])
27
28 # Concatenate all feature maps along the patch dimension
29 memory_bank = torch.cat(memory_bank, dim=0).cuda() # Shape: (total_patches_across_all_images, feature_dim)
30
31 # Output the final memory bank shape
32 memory_bank.shape # e.g., torch.Size([306544, 391]) = 784 patches x num_images, each of 391 features
33

```

Downloading: "<https://download.pytorch.org/models/resnet50-11ad3fa6.pth>" to /root/.cache/torch/hub/checkpoints/resnet50-11ad3fa6.pth [100%] 97.8M/97.8M [00:00<00:00, 207MB/s]

```

220 torch.Size([784, 1536])
221 torch.Size([172480, 1536])

```

```

1 # Randomly select 10% of the feature vectors from the memory bank without replacement
2 selected_indices = np.random.choice(len(memory_bank), size=len(memory_bank)//10, replace=False)
3
4 # Subsample the memory bank using the selected indices
5 memory_bank = memory_bank[selected_indices]
6
7 # Output the new shape of the memory bank (should be 10% of the original size)
8 memory_bank.shape
9

```

torch.Size([17248, 1536])

```

1 # Initialize a list to store the anomaly scores (y_score_good) for all "good" training images
2 y_score_good = []
3
4 # Define the folder path again to iterate over all "good" images from the training set
5 folder_path = Path('/content/drive/MyDrive/mvttec_anomaly_detection/metal_nut/train/good')
6
7 # Iterate through each image file in the folder
8 for pth in tqdm(folder_path.iterdir(), leave=False):
9     # Load and preprocess the image: resize, normalize, and add batch dimension
10     data = transform(Image.open(pth).convert("RGB")).cuda().unsqueeze(0) # Shape: (1, 3, 224, 224)

```

```

11
12     with torch.no_grad(): # Disable gradient tracking for inference
13         features = backbone(data) # Extract patch features → shape: (784, 391)
14
15     # Compute pairwise Euclidean distances between extracted features and memory bank
16     distances = torch.cdist(features, memory_bank, p=2.0) # Shape: (784, memory_bank_size)
17
18     # For each patch in the image, find the minimum distance to the memory bank
19     dist_score, dist_score_idxs = torch.min(distances, dim=1) # dist_score: (784,)
20
21     # Use the maximum of all minimum distances as the anomaly score for the image
22     s_star = torch.max(dist_score) # Higher score → more anomalous
23
24     # Optionally reshape per-patch distance scores into a 28x28 segmentation map
25     segm_map = dist_score.view(1, 1, 28, 28) # Useful for pixel-level anomaly visualization
26
27     # Store the image-level anomaly score
28     y_score_good.append(s_star.cpu().numpy())
29     # break # Uncomment for debugging a single iteration
30
31 # Preview the first 5 image-level anomaly scores
32 y_score_good[:5]
33

```

```

↗ [array(13.765262, dtype=float32),
   array(12.175159, dtype=float32),
   array(14.22541, dtype=float32),
   array(12.460357, dtype=float32),
   array(14.071746, dtype=float32)]


```

```

1 # Preview the first 5 image-level anomaly scores for "good" training images
2 y_score_good[:5]
3
4 # Remove batch and channel dimensions from the segmentation map
5 # Original shape: (1, 1, 28, 28) → After squeeze: (28, 28)
6 image_np = segm_map.squeeze().cpu()
7
8 # Plot the 28x28 anomaly score heatmap (distance map)
9 plt.imshow(image_np, cmap='gray') # Use grayscale colormap
10 plt.title("28x28 Image") # Title of the plot
11 plt.axis("off") # Hide x and y axis for cleaner view
12 plt.show()
13

```

↗ 28x28 Image



```

1 # Calculate the mean anomaly score of all "good" training images
2 print(np.mean(y_score_good)) # Expected to be relatively low
3
4 # Calculate the standard deviation of anomaly scores
5 print(np.std(y_score_good)) # Indicates spread/variance in anomaly scores
6
7 # Define the anomaly detection threshold as  $\mu + 3\sigma$  (3-sigma rule)
8 # Assumes "good" scores follow a roughly Gaussian distribution
9 best_threshold = np.mean(y_score_good) + 3 * np.std(y_score_good)


```

```

10 print(f"Threshold: {best_threshold}")
11
12 # Plot histogram of anomaly scores for "good" images
13 plt.hist(y_score_good, bins=50)
14
15 # Draw a vertical red line at the computed threshold
16 plt.vlines(x=best_threshold, ymin=0, ymax=30, color='r')
17
18 # Show the plot
19 plt.show()
20

```

↔ 13.309807  
0.73102367  
Threshold: 15.502878189086914



```

1 # Lists to store predicted anomaly scores and ground truth labels
2 y_score = [] # Image-level anomaly scores
3 y_true = [] # Corresponding ground truth labels (0 = good, 1 = defect)
4
5 # Path to test dataset root folder
6 base_test_path = Path("/content/drive/MyDrive/mvtec_anomaly_detection/metal_nut/test")
7
8 # Iterate over all defect types (including 'good')
9 for defect_type in ['bent', 'color', 'flip', 'good', 'scratch']:
10     folder_path_test = base_test_path / defect_type # Path to current subfolder
11
12     for pth in tqdm(folder_path_test.iterdir(), leave=False):
13         class_label = pth.parts[-2] # Get defect category from path (e.g., 'good', 'rough', etc.)
14
15         with torch.no_grad(): # Disable gradient computation
16             # Load and preprocess image
17             test_image = transform(Image.open(pth).convert("RGB")).cuda().unsqueeze(0)
18
19             # Extract features using the backbone model
20             features = backbone(test_image) # Shape: (784, 391)
21
22             # Compute L2 distance between extracted features and memory bank
23             distances = torch.cdist(features, memory_bank, p=2.0)
24
25             # For each patch, find the minimum distance to memory bank
26             dist_score, _ = torch.min(distances, dim=1)
27
28             # Maximum patch distance is the image-level anomaly score
29             s_star = torch.max(dist_score)
30
31             # Optional: create 28x28 segmentation map from patch distances (useful for heatmap visualization)
32             segm_map = dist_score.view(1, 1, 28, 28)
33
34             # Append results
35             y_score.append(s_star.cpu().numpy())
36             y_true.append(0 if class_label == 'good' else 1) # 0 = normal, 1 = anomaly
37
38

```




```
1 y_score[40:45], y_true[40:45]
```



```
(array(23.366669, dtype=float32),
 array(20.992207, dtype=float32),
 array(14.004607, dtype=float32),
 array(18.815985, dtype=float32),
 array(17.973526, dtype=float32)),
 [1, 1, 1, 1, 1])
```

```
1 # Filter out anomaly scores (y_score) that correspond to the 'BAD' (defective) class
2 y_score_bad = [score for score, true in zip(y_score, y_true) if true == 1]
3
4 # Plot a histogram of anomaly scores for BAD samples
5 plt.hist(y_score_bad, bins=50)
6
7 # Draw a vertical red line at the threshold (computed earlier using 3-sigma rule)
8 plt.vlines(x=best_threshold, ymin=0, ymax=30, color='r')
9
10 # Display the plot
11 plt.show()
12
```



```
1 # Path to a specific defective test image
2 test_image_path = '/content/drive/MyDrive/mvtec_anomaly_detection/metal_nut/test/scratch/000.png'
3
4 # Extract features from the input image using the pretrained backbone
5 features = backbone(test_image) # test_image must be defined before this line
6
7 # Compute pairwise Euclidean distances between patches and memory bank
8 distances = torch.cdist(features, memory_bank, p=2.0)
9
10 # For each patch, find the closest (minimum distance) memory bank feature
11 dist_score, dist_score_idx = torch.min(distances, dim=1)
12
13 # Use the maximum patch-level distance as the image-level anomaly score
14 s_star = torch.max(dist_score)
15
16 # Reshape per-patch distances into a 28x28 segmentation map
17 segm_map = dist_score.view(1, 1, 28, 28)
18
19 # Upscale the segmentation map to 224x224 to match the original image resolution
20 # This makes it visually align with the input image if you overlay it
21 segm_map = torch.nn.functional.interpolate(
22     segm_map,
23     size=(224, 224),
24     mode='bilinear'
25 )
26
27 # Plot the upscaled anomaly heatmap using 'jet' colormap for better visual contrast
28 plt.figure(figsize=(4, 4))
29 plt.imshow(segm_map.cpu().squeeze(), cmap='jet')
```


```

30 plt.title("Anomaly Heatmap")
31 plt.axis("off")
32 plt.show()
33

```



Anomaly Heatmap



```

1 from sklearn.metrics import roc_auc_score, roc_curve, confusion_matrix, ConfusionMatrixDisplay, f1_score
2
3 # -----
4 # Step 1: Compute AUC-ROC score
5 # -----
6 # Measures the model's ability to distinguish between GOOD (0) and BAD (1) images
7 auc_roc_score = roc_auc_score(y_true, y_score)
8 print("AUC-ROC Score:", auc_roc_score)
9
10 # -----
11 # Step 2: Generate ROC curve
12 # -----
13 # Returns False Positive Rate, True Positive Rate, and thresholds
14 fpr, tpr, thresholds = roc_curve(y_true, y_score)
15 print("fpr, tpr, thresholds: ", fpr, tpr, thresholds)
16
17 # -----
18 # Step 3: Calculate F1 score for each threshold
19 # -----
20 # Evaluate which threshold gives the best balance of precision and recall
21 f1_scores = [f1_score(y_true, y_score >= threshold) for threshold in thresholds]
22 print("f1_scores:", f1_scores)
23
24 # Select threshold that yields the highest F1 score
25 best_threshold = thresholds[np.argmax(f1_scores)]
26 print(f'best_threshold = {best_threshold}')
27
28 # -----
29 # Step 4: Plot ROC curve
30 # -----
31 plt.figure()
32 plt.plot(fpr, tpr, color='darkorange', lw=2,
33          label='ROC curve (area = %0.2f)' % auc_roc_score)
34 plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--') # Diagonal line = random guess
35 plt.xlabel('False Positive Rate')
36 plt.ylabel('True Positive Rate')
37 plt.title('Receiver Operating Characteristic (ROC) Curve')
38 plt.legend(loc="lower right")
39 plt.show()
40
41 # -----
42 # Step 5: Display Confusion Matrix
43 # -----
44 # Use the best threshold to classify predictions
45 cm = confusion_matrix(y_true, (y_score >= best_threshold).astype(int))
46 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['GOOD', 'BAD'])
47 disp.plot()
48 plt.title("Confusion Matrix at Best Threshold")
49 plt.show()
50


```

```


AUC-ROC Score: 0.9853372434017595
fpr, tpr, thresholds: [0. 0. 0. 0.04545455 0.04545455 0.09090909
0.09090909 0.18181818 0.18181818 0.54545455 0.54545455 1.
] [0. 0.01075269 0.90322581 0.90322581 0.9462365
0.95698925 0.95698925 0.98924731 0.98924731 1.
] [
14.918959 14.770254 14.188081 14.004607 12.644429]
f1_scores: [0.0, 0.02127659574468085, 0.9491525423728814, 0.9438202247191011, 0.967032967032967, 0.9617486338797814, 0.967
best_threshold = 14.770254135131836

```

Receiver Operating Characteristic (ROC) Curve



Confusion Matrix at Best Threshold



```

1 import os
2 import cv2
3 import time
4 import torch
5 import matplotlib.pyplot as plt
6 from pathlib import Path
7 from PIL import Image
8
9 backbone.eval()
10 class_label = ['GOOD', 'BAD']
11 wanted_types = {'bent', 'color', 'flip', 'good', 'scratch'}
12
13 test_path = Path('/content/drive/MyDrive/mvtec_anomaly_detection/metal_nut/test')
14 output_root = Path('/content/outputs')
15
16 # Track which fault types have already been displayed
17 displayed_faults = set()
18
19 for path in test_path.glob('/*.png'):
20     fault_type = path.parts[-2]

```

```

21 item_name = path.parts[-4]
22
23 if fault_type not in wanted_types:
24     continue
25
26 # Load and transform the image
27 test_image = transform(Image.open(path).convert("RGB")).cuda().unsqueeze(0)
28
29 with torch.no_grad():
30     features = backbone(test_image)
31
32 distances = torch.cdist(features, memory_bank, p=2.0)
33 dist_score, _ = torch.min(distances, dim=1)
34 s_star = torch.max(dist_score)
35
36 segm_map = dist_score.view(1, 1, 28, 28)
37 segm_map = torch.nn.functional.interpolate(
38     segm_map, size=(224, 224), mode='bilinear'
39 ).cpu().squeeze().numpy()
40
41 y_score_image = s_star.cpu().numpy()
42 y_pred_image = 1 * (y_score_image >= best_threshold)
43
44 # Create output path
45 save_path = output_root / item_name / 'test' / fault_type
46 save_path.mkdir(parents=True, exist_ok=True)
47
48 # Plot
49 fig, axs = plt.subplots(1, 3, figsize=(12, 3))
50
51 axs[0].imshow(test_image.squeeze().permute(1, 2, 0).cpu().numpy())
52 axs[0].set_title(f'Fault Type: {fault_type}')
53 axs[0].axis('off')
54
55 axs[1].imshow(segm_map, cmap='jet', vmin=best_threshold, vmax=best_threshold * 2)
56 axs[1].set_title(f'Score: {y_score_image:.2f} | {class_label[y_pred_image]}')
57 axs[1].axis('off')
58
59 axs[2].imshow((segm_map > best_threshold), cmap='gray')
60 axs[2].set_title('Segmentation Map')
61 axs[2].axis('off')
62
63 plt.tight_layout()
64
65 # Save visualization
66 out_file = save_path / f'{path.stem}_vis.png'
67 fig.savefig(out_file)
68
69 # Display only the first image per fault type
70 if fault_type not in displayed_faults:
71     plt.show()
72     displayed_faults.add(fault_type)
73 else:
74     plt.close(fig)
75

```



